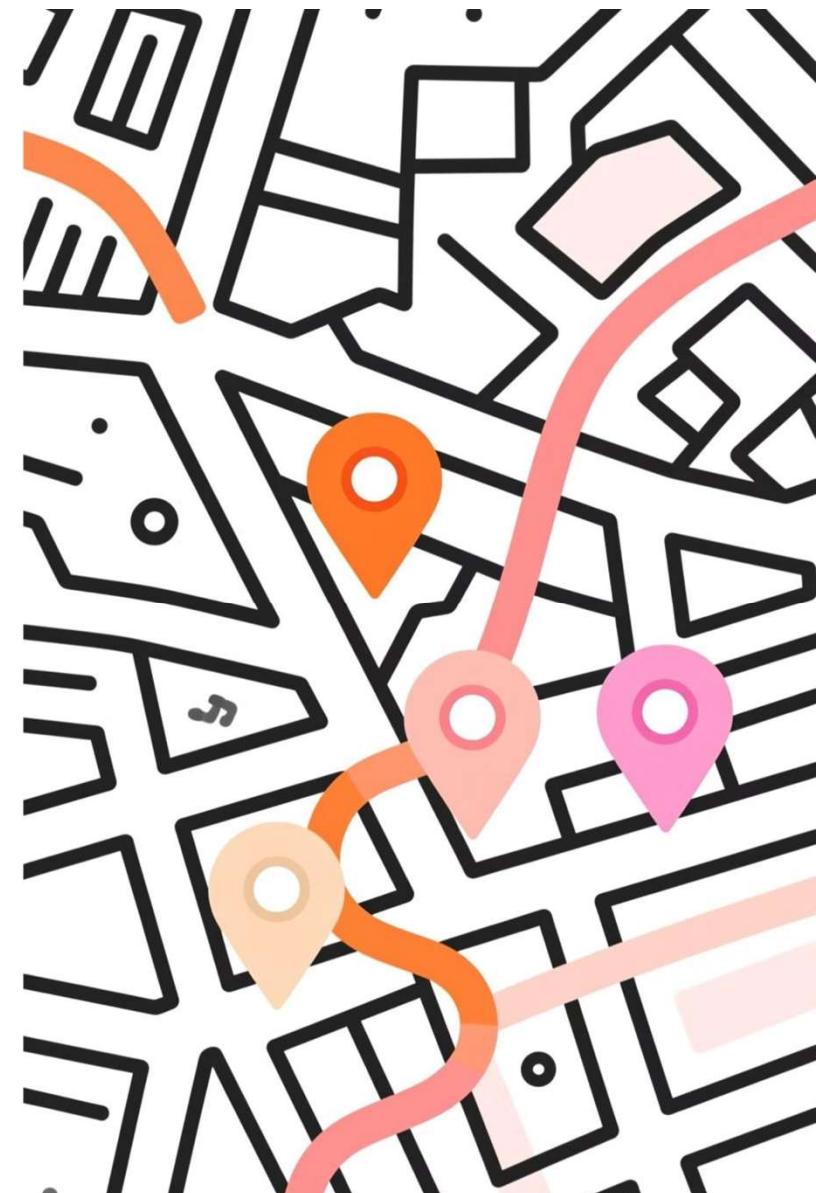
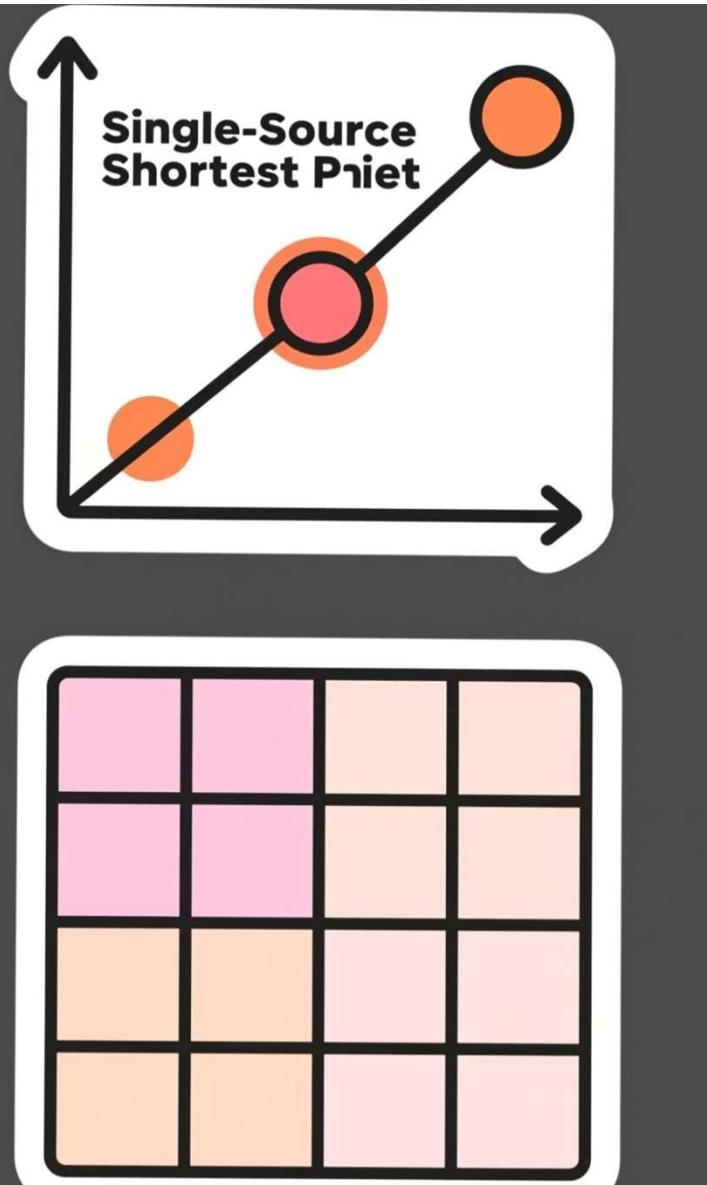


Shortest Paths: Dijkstra vs. Floyd–Warshall

This presentation compares two foundational algorithms for shortest paths in graphs: Dijkstra (single-source shortest paths, SSSP) and Floyd–Warshall (all-pairs shortest paths, APSP). We'll cover core ideas, algorithmic steps, complexity trade-offs, practical examples, and a performance comparison to guide when to use each method.





Problem definitions: SSSP vs. APSP

Single-Source Shortest Path (SSSP)

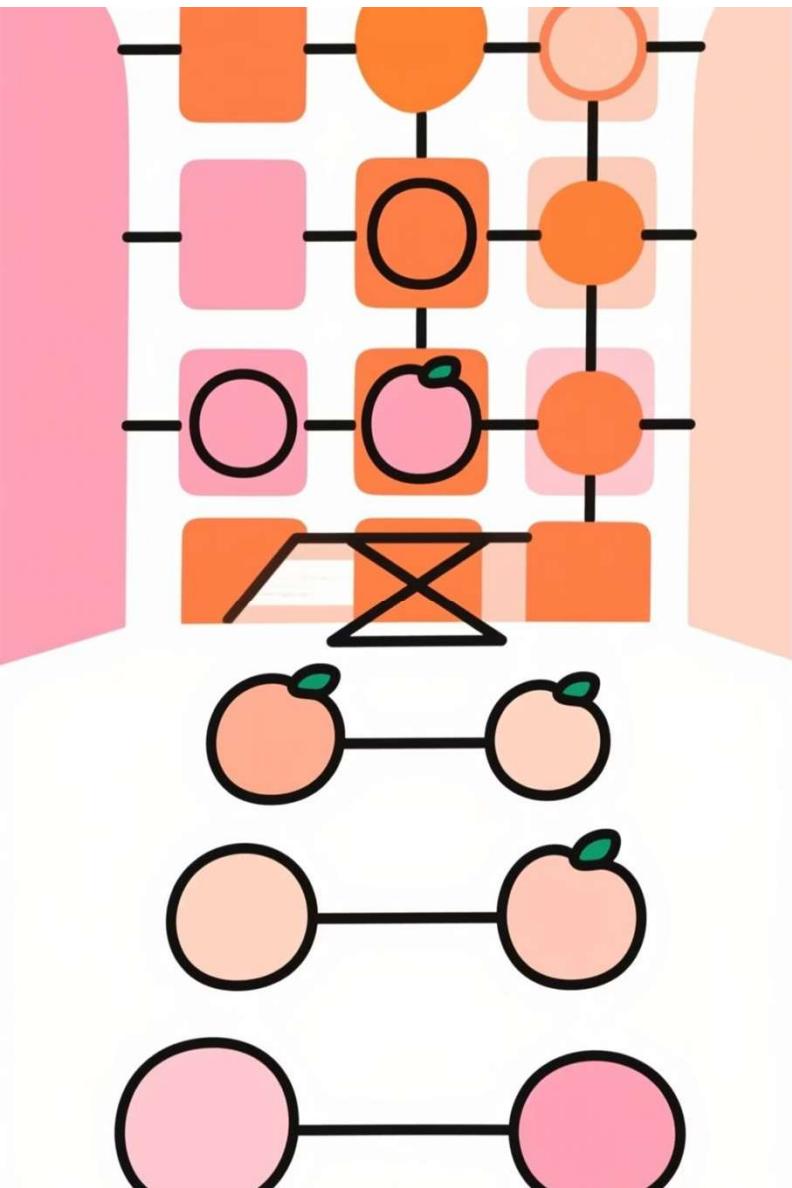
Given a weighted graph $G = (V, E)$ and source s , compute shortest path distances from s to every other node. Typical use: routing from one origin to many destinations.

- One source, many targets
- Often uses Dijkstra (non-negative weights)
- Complexity depends on $|V|$, $|E|$ and data structure

All-Pairs Shortest Paths (APSP)

Compute shortest path distances between every pair of nodes (u, v) in the graph. Useful when queries are many or when you need global distance information.

- Every source \rightarrow every destination
- Floyd-Warshall is a classical solution
- Costs scale with $|V|^3$ for dense approaches



How Dijkstra's Algorithm Works

Dijkstra is a greedy algorithm for SSSP on graphs with non-negative edge weights. Initialize distances ($\text{dist}[s]=0$, others = ∞). Repeatedly extract the vertex with smallest tentative distance from a min-priority queue, then relax its outgoing edges (update neighbor distances if a shorter path is found). Once extracted, a vertex's distance is finalized.

Initialize

Set $\text{dist}[s]=0$, push s into the priority queue.

Extract-min

Pop vertex u with smallest $\text{dist}[u]$; u is now settled.

Relax edges

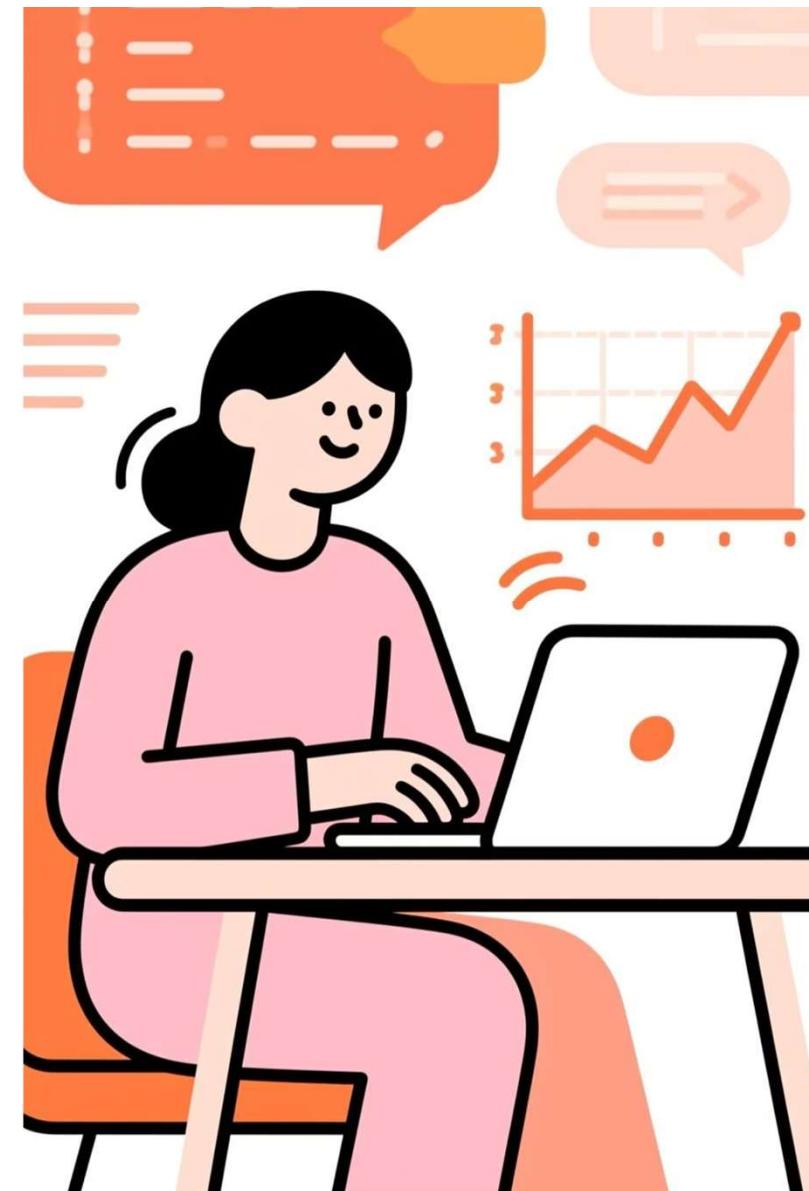
For each (u,v,w) : if $\text{dist}[u]+w < \text{dist}[v]$, update $\text{dist}[v]$ and decrease-key.

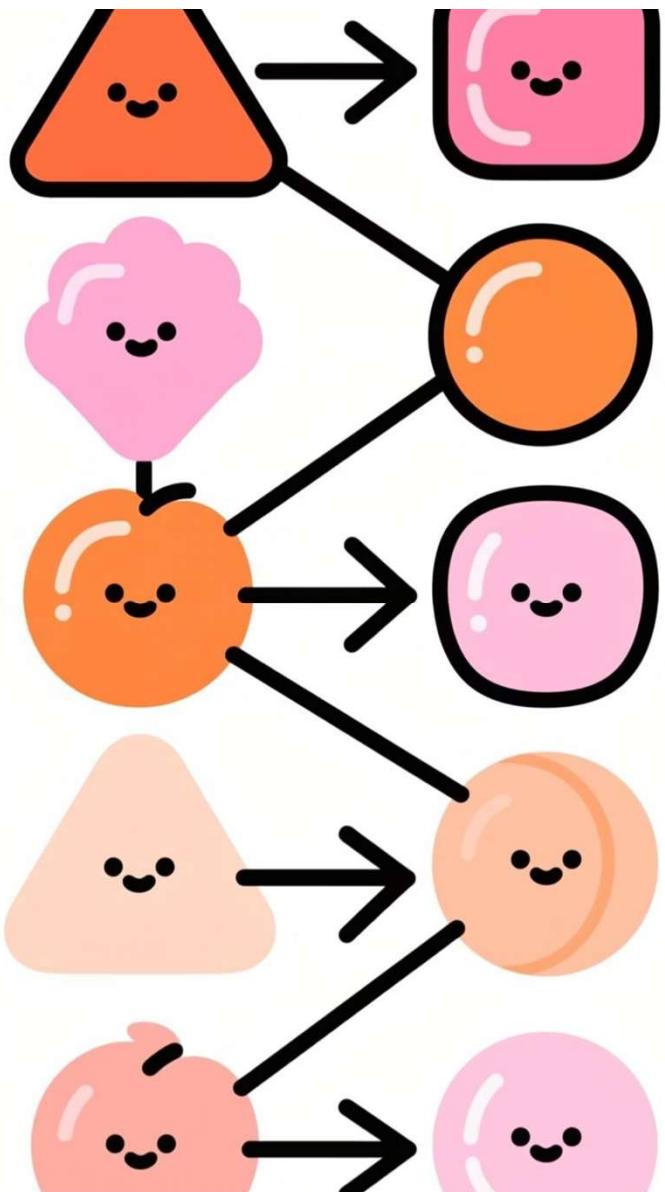
Dijkstra: Complexity & Limitations

Time complexity depends on implementation: using a binary heap: $O((|V|+|E|) \log |V|) \approx O(|E| \log |V|)$ for sparse graphs. using a Fibonacci heap: $O(|E| + |V| \log |V|)$. Space complexity: $O(|V|)$ for distances and $O(|V|+|E|)$ for graph storage.

- Fast for sparse graphs with single-source queries
- Requires non-negative edge weights
- Repeated runs needed for many different sources

Note: Dijkstra cannot handle negative-weight edges – Bellman-Ford or Johnson's algorithm are alternatives when negative weights exist.





How Floyd–Warshall Works

Floyd–Warshall computes APSP using dynamic programming on an adjacency matrix. Let $\text{dist}[i][j]$ be current best distance. Iterate k from $1..n$ and allow paths that can use only vertices $\{1..k\}$ as intermediates. Update rule: $\text{dist}[i][j] = \min(\text{dist}[i][j], \text{dist}[i][k] + \text{dist}[k][j])$. After n iterations, dist contains shortest distances for all pairs.



Initialize matrix

$\text{dist}[i][j] = \text{weight}(i,j)$ or ∞ if no edge;
 $\text{dist}[i][i]=0$.



Iterate k

For each pair (i,j) check if path through k improves $\text{dist}[i][j]$.



Result

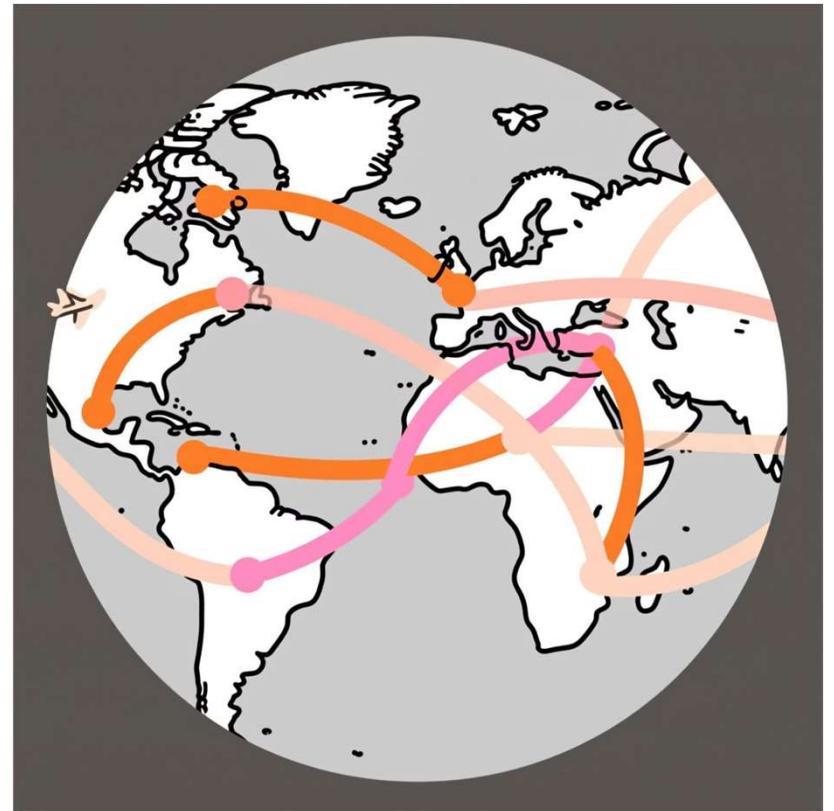
Matrix holds shortest distances for every (i,j) .

Floyd–Warshall: Complexity & Benefits

Time complexity: $O(|V|^3)$ regardless of sparsity – predictable but expensive for large graphs. Space complexity: $O(|V|^2)$ to store the distance matrix (and optionally next-hop reconstruction table).

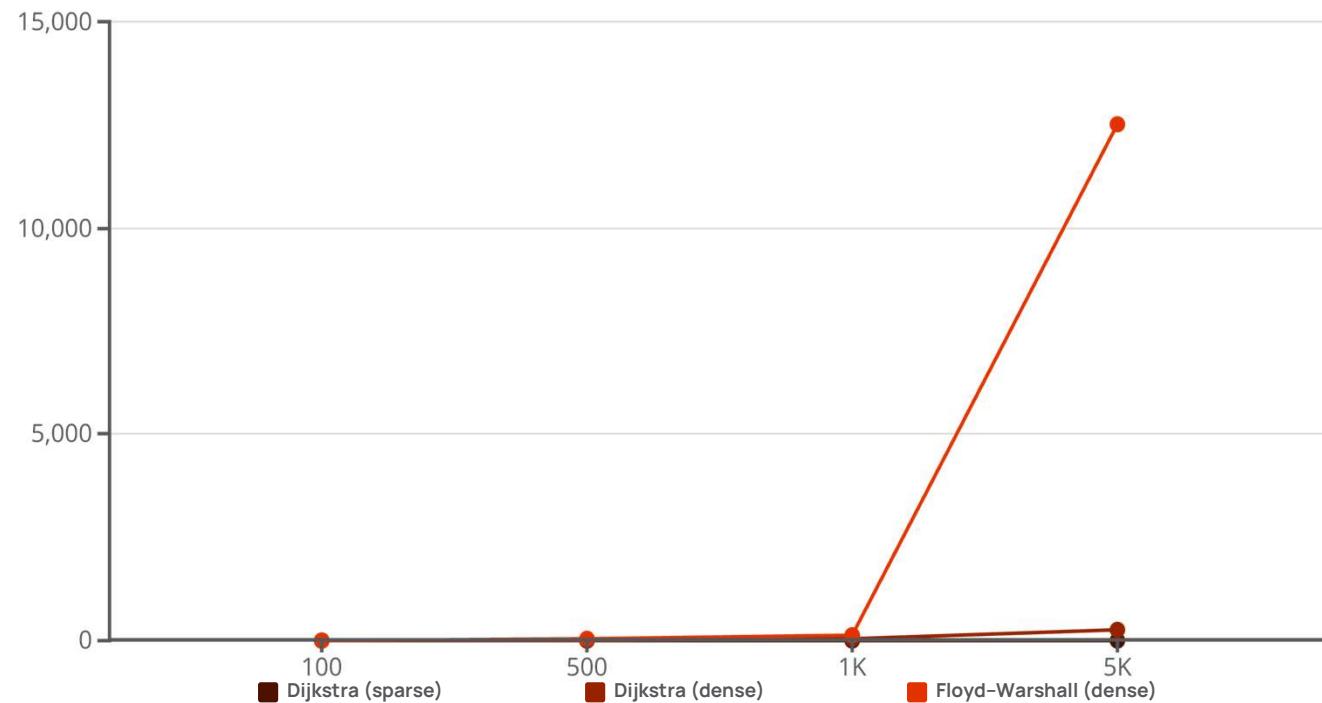
Strengths: handles negative edge weights (detects negative cycles), simple matrix-oriented implementation, natural fit for dense graphs and problems needing distances between all pairs.

- APSP computed in one run
- Detects negative cycles via $\text{dist}[i][i] < 0$
- Easy to implement and parallelize (matrix operations)



Performance Comparison: Time vs. Vertices

Comparison setup: synthetic undirected graphs with varying $|V|$ and two regimes:
- Sparse: $|E| \approx c|V|$ (c small)
- Dense: $|E| \approx |V|(|V|-1)/2$
We'll compare single-run runtimes for Dijkstra (binary heap) and Floyd-Warshall.



Interpretation: Dijkstra scales well for sparse graphs and single-source queries. Floyd-Warshall is competitive only for small-to-moderate vertex counts or when you require APSP and the graph is dense.

Real-world Example: Google Maps & Dijkstra

Google Maps uses variants of SSSP (A*, Dijkstra-like heuristics) to compute fastest route from current location to a destination. Characteristics that make Dijkstra-based approaches suitable:



Single-source queries

Each navigation request has one origin and one (or a few) destinations.



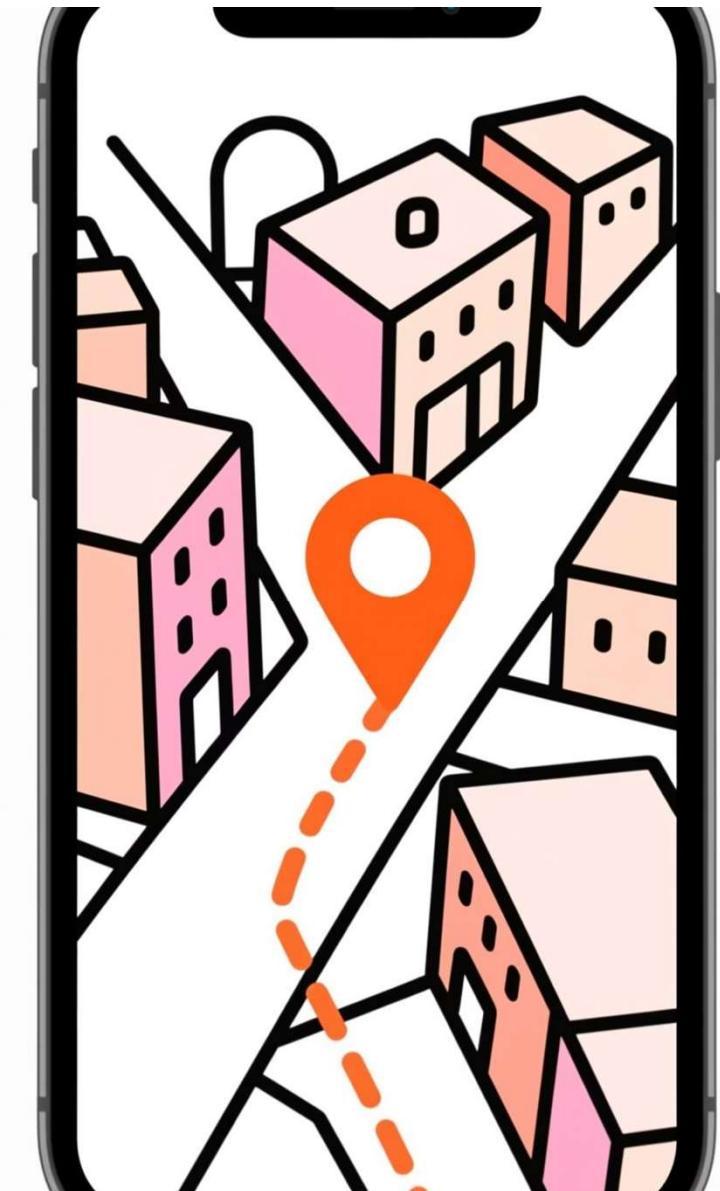
Sparse road networks

Road graphs are sparse: IEI is $O(|V|)$, making Dijkstra fast with proper heuristics.



Heuristics & constraints

A* adds admissible heuristics (straight-line distance) to speed up search for large maps.



Real-world Example: Airline Route Planning & Floyd–Warshall

Airline route planning often requires global connectivity and frequent APSP queries: e.g., computing minimum hop or minimum-cost connections between any pair of airports, or preparing lookup tables for scheduling and delay impact analysis.



Precomputation advantage

When operations need all-pairs distances readily available, Floyd–Warshall produces a complete distance matrix in one run.

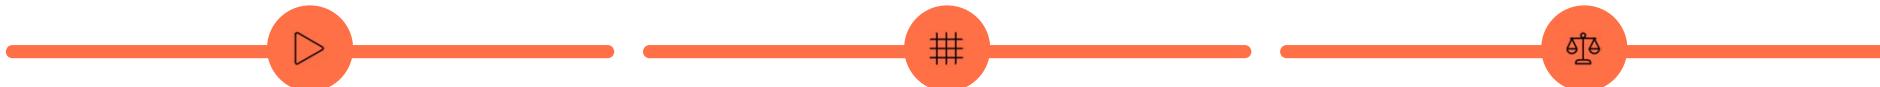


Dense connectivity

Hub-and-spoke networks with dense effective connectivity (many indirect routes) make matrix-based approaches more plausible for moderate IVI.

Choosing Between Dijkstra and Floyd–Warshall

Summary guidance for algorithm selection:



Use Dijkstra (SSSP)

When you have one source (or a few) and a large sparse graph, need fast queries per source, and edges are non-negative. Combine with A* or goal-directed techniques for large spatial graphs.

Use Floyd–Warshall (APSP)

When you need distances between all pairs, graph size is moderate ($|V|$ up to a few thousand depending on resources), or you need negative-weight support and negative-cycle detection.



Hybrid & Practical Tips

For many-sources but not all, consider repeated Dijkstra runs, Johnson's algorithm (reweight + Dijkstra) for sparse graphs with negative edges, or precompute APSP selectively for important nodes (landmarks).

Final takeaway: match the algorithm to the query pattern and graph structure. Use Dijkstra-based approaches for targeted routing at scale; use Floyd–Warshall for complete global distance tables or when negative weights and simplicity matter.